

Objective

- Practice using the Standard Library containers, iterators, and algorithms.
1. Recall that a palindrome is a word or phrase that reads the same when read forward or backward, such as "Was it a car or a cat I saw?". The reading process ignores spaces, punctuation, and capitalization.

Write a function named `isPalindrome` that receives a string as parameter and determines whether the received string is a palindrome.

Your implementation may *not* use

- any loops explicitly; that is, no `for`, `while` or `do/while` loops
- more than one local `string` object
- raw arrays, STL container classes

Use the following function to test your `isPalindrome` function:

```
void test_is_palindrome()
{
    std::string str_i = std::string("Was it a car or a cat I saw?");
    std::string str_u = std::string("Was it a car or a cat U saw?");
    cout << "the phrase \"" + str_i + "\"" is " +
        (is_palindrome(str_i) ? "" : "not ") + "a palindrome\n";
    cout << "the phrase \"" + str_u + "\"" is " +
        (is_palindrome(str_u) ? "" : "not ") + "a palindrome\n";
}
```

A Suggestion:

```
bool is_palindrome(const std::string & phrase)
{
    string temp;

    1. use std::remove_copy_if to move only alphabet characters from phrase to
       temp; note that temp is empty; hence the need for an inserter iterator
    2. use std::transform to convert all characters in temp to uppercase (or
       lowercase) (see page 22, Smart Pointers + Move Semantics, Moodle)
    3. use std::equal to compare the first half of temp with its second half,
       moving forward in the first half (starting at temp.begin()) and moving
       backward in the second half (starting at temp.rbegin())
    4. return the outcome of step 3
}
```

2. Write a function template named `second_max` to find the second largest element in a container within a given range `[start, finish)`, where `start` and `finish` are iterators that provide properties of forward iterators.

Your function template should be prototyped as follows, and may not use STL algorithms or containers.

```
template <class Iterator> // template header
std::pair<Iterator,bool> // function template return type
second_max(Iterator start, Iterator finish) // function signature
{
    // your code
}
```

Clearly, in the case where the iterator range `[start, finish)` contains at least two distinct objects, `second_max` should return an iterator to the second largest object. However, what should `second_max` return if the iterator range `[start, finish)` is empty or contains objects which are all equal to one another? How should it convey all that information back to the caller?

Mimicking `std::set's` `insert` member function, your `second_max` function should return a `std::pair<Iterator,bool>` defined as follows:

condition	value to return
R is empty	<code>std::make_pair (finish,false)</code>
R contains all equal elements	<code>std::make_pair (start,false)</code>
R contains at least two distinct elements	<code>std::make_pair (iter,true)</code>

R is the range `[start, finish)`,
iter is an `Iterator` referring to the 2nd largest element in the range.

Use the following function to test your `second_max` function:

```

void test_second_max(std::vector<int> vec)
{
    // note: auto in the following statement is deduced as
    // std::pair<std::vector<int>::iterator, bool>
    auto retval = second_max(vec.begin(), vec.end());

    if (retval.second)
    {
        cout << "The second largest element in vec is "
              << *retval.first << endl;
    }
    else
    {
        if (retval.first == vec.end())
            cout << "List empty, no elements\n";
        else
            cout << "Container's elements are all equal to "
                  << *retval.first << endl;
    }
}

```

3. Consider the `count_if` algorithm in the `<algorithm>` header:

```
function template std::count_if
```

```
template <class InputIterator, class UnaryPredicate>  
    typename iterator_traits<InputIterator>::difference_type  
    count_if (InputIterator first, InputIterator last, UnaryPredicate pred);
```

Returns the number of elements in the range `[first,last)` for which `pred` is true.

Write a function named `countStrings` that takes two parameters: a `vector<std::string>` (by reference), and an integer `n`. Your function must return the number of strings in the vector whose length is equal to `n`. Your functions must use the `count_if` algorithm.

For example, suppose

```
std::vector<std::string> vec { "C", "BB", "A", "CC", "A", "B",  
                              "BB", "A", "D", "CC", "DDD", "AAA" };
```

Then, for example, the call `countStrings(vec, 1)` should return 6, the call `countStrings(vec, 3)` should return 2, etc.

You should write three versions of `countStrings`, implementing the callable argument to `count_if` as follows:

- A. A lambda expression named `countStringLambda`
- B. A free function named `countStringFreeFun`
- C. A functor (function object) named `countStringFunctor`

Recall that an object or expression is callable if the call operator can be applied to it.

4. The STL algorithm `for_each` accepts as parameters a range of iterators and a unary callable, then calls the callable on each argument.

Write a function `charFrequency()` that reads a string of any number of characters from the keyboard, storing and counting the frequencies of each input character (except whitespace).

Your code must use a `map<char, int>` container to store the characters (keys) and their frequency counts (values) as `pair<char, int>` objects. Do not use any loops; instead, use the `for_each` algorithm to extract characters from the input stream and reflect the characters and their frequency counts into the map.

Hint: although the task at hand here might seem very involved, reading input characters and counting their frequencies can be implemented using just one statement! Simply, implement the callable parameter of the `for_each` algorithm as a lambda expression, capturing your map object into your lambda by reference.

Here is the output from a sample run:

```
Enter one or more lines of text.
To end input press Ctrl-Z in Windows and Ctrl-D in Linux
line one
line two
line three and no more lines after this last line.
^Z
. 1
a 3
d 1
e 10
f 1
h 2
i 6
l 6
m 1
n 8
o 4
r 3
s 3
t 5
w 1
```

5. Consider the following function that uses a `multiset` defined using `std::multiset`'s default compare function:

```
1 void multisetUsingDefaultComparator()
2 {
3     std::multiset<std::string> strSet; // an empty set
4
5     // a set that uses the default std::less<int> to sort the set elements
6     std::vector<std::string> vec {"C", "BB", "A", "CC", "A", "B",
7                                   "BB", "A", "D", "CC", "DDD", "AAA" };
8
9     // copy the vector elements to our set.
10    // We must use a general (not front or back) inserter
11    // (set does not have push_front or push_back members,
12    // so we can't use front or back inserter)
13    std::copy(vec.begin(), vec.end(), // source container range
14              std::inserter(strSet, strSet.begin())); // general inserter
15
16    // create an ostream_iterator for writing to cout,
17    // using a space " " as a separator
18    std::ostream_iterator<std::string> out(cout, " ");
19
20    // output the set elements to cout separating them with a space
21    std::copy(strSet.begin(), strSet.end(), out);
22 }
```

The function produces the following output:

```
A A A AAA B BB BB C CC CC D DDD
```

Modify the construction of the `strSet` set at line 3 so that the resulting function, named `multisetUsingMyComparator()`, outputs:

```
A A A B C D BB BB CC CC AAA DDD
```

The effect is that the string elements in `strSet` are now ordered into groups of strings of increasing lengths 1, 2, 3, ..., with the strings in each group sorted lexicographically.

Test Driver Code

```
int main()
{
    // problem 1:
    test_is_palindrome();
    cout << "\n";

    // problem 2:
    std::vector<int> v1{ 1 }; // one element
    test_second_max(v1);

    std::vector<int> v2{ 1, 1 }; // all elements equal
    test_second_max(v2);

    std::vector<int> v3{ 1, 1, 3, 3, 7, 7 }; // at least with two distinct elements
    test_second_max(v3);
    cout << "\n";

    // problem 3:
    std::vector<std::string> vecstr
    { "count_if", "Returns", "the", "number", "of", "elements", "in", "the",
      "range", "[first", "last)", "for", "which", "pred", "is", "true."
    };
    cout << countStringLambda(vecstr, 5) << endl;
    cout << countStringFreeFun(vecstr, 5) << endl;
    cout << countStringFunctor(vecstr, 5) << endl;
    cout << "\n";

    // problem 4:
    charFrequency();
    cout << "\n";

    // problem 5:
    multisetUsingMyComparator();
    cout << "\n";

    return 0;
}
```

1 Deliverables

Implementation files: Only one file named `A5_all.cpp` that includes `is_palindrome`, `test_is_palindrome`, `second_max`, `test_second_max`, `countStringLambda`, `countStringFreeFun`, `countStringFunctor`, `charFrequency`, `multisetUsingMyComparator`, as well as all the supporting functions, functors, and lambdas.

README.txt A text file, as described in the course outline.

2 Marking scheme

	Program correctness:
80%	16% per problem
5%	Format, clarity, completeness of output
10%	Javadoc style documentation before introduction of every class and function, Concise documentation of nontrivial steps in code, choice of variable names, indentation and readability of program